

Maximizing Throughput in Distributed Systems: Java & Go Servers, Java Client

Random Team: Yinuo Feng, Nuanxin Jin, Boyuan Sun, Zihe Chen, Chenyu Jiang

Table of contents

01

Introduction

02

The Problem

03

**Proposed
System
Architecture**

04

**Experiments
& Results**

05

Analysis

06

Future Works

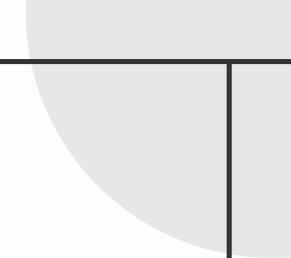
Tech Stack

Servers: Java & Go

Http Client: Java

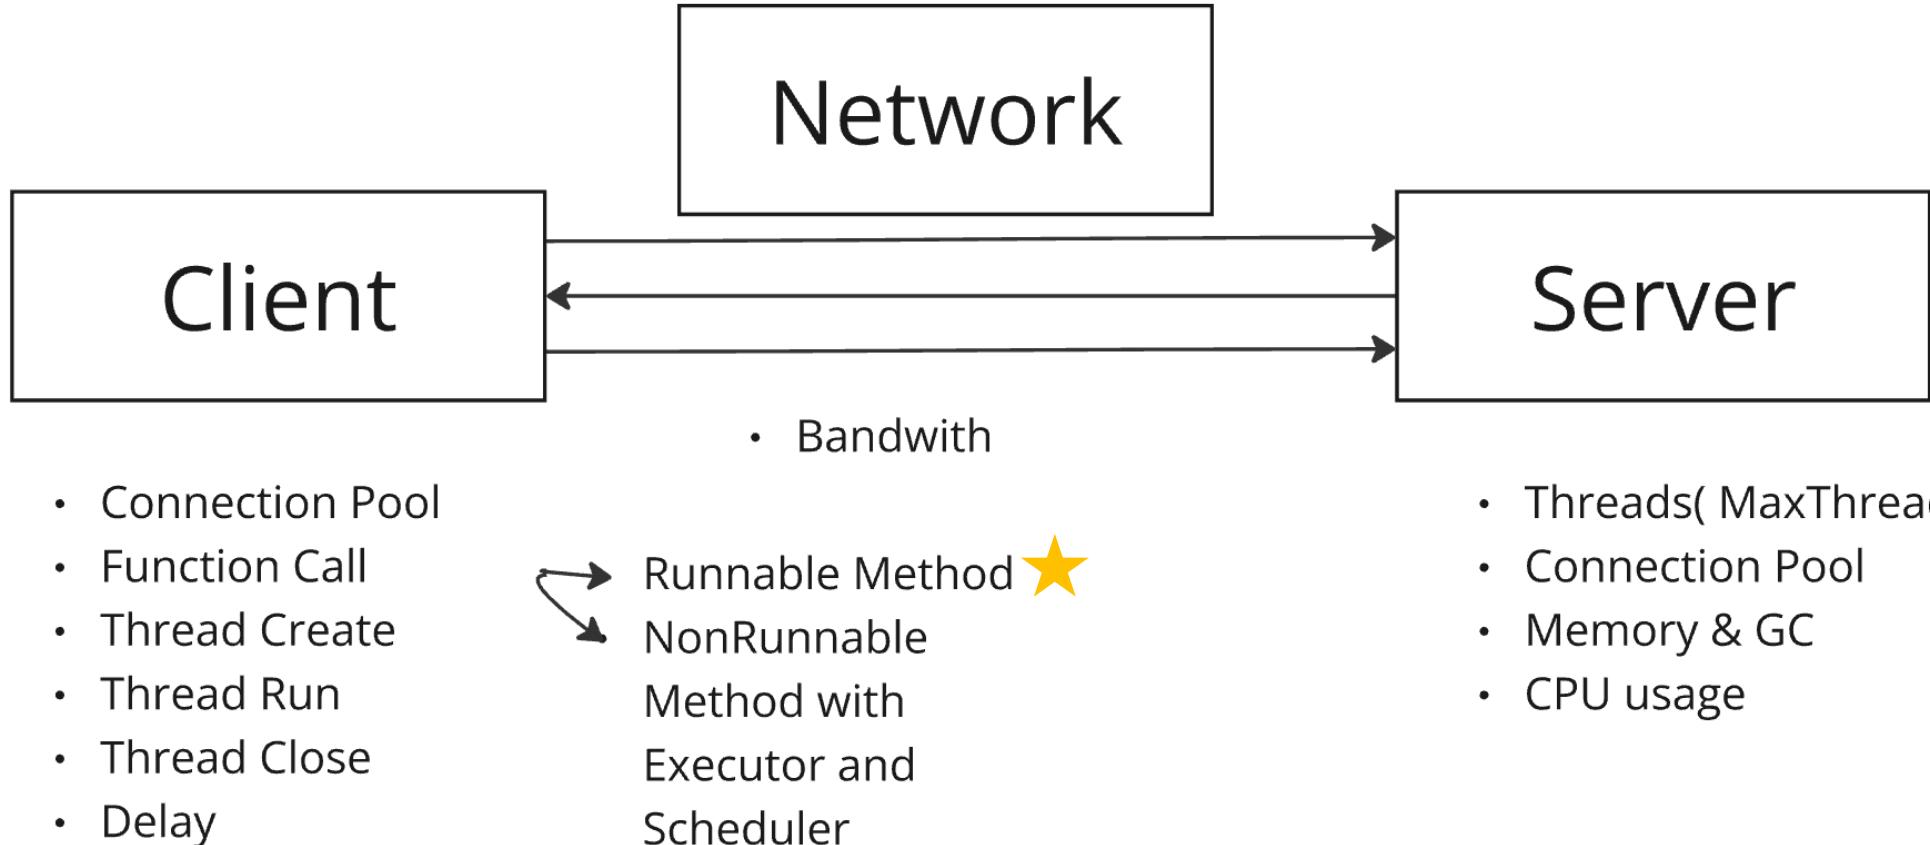
Infrastructure: AWS EC2, Tomcat

Testing & Metrics: Wall time, Throughput, Latency, p99 response time



How to maximize throughput?

Proposed System Architecture



Runnable Method

- Runnable Interface executed by Thread

```
private CloseableHttpClient client; 2 usages
private List<Row> data; 2 usages

public ClientGet(String IPAddr, CloseableHttpClient client, List<Row> data) {
    this.getUrl = "http://" + IPAddr + "/IGORTON/AlbumStore/1.0.0/albums";
    this.client = client;
    this.data = data;
}

// stolen from https://hc.apache.org/httpclient-legacy/tutorial.html
public void run() {

    // Create a method instance.
    CountDownLatch completed2 = new CountDownLatch(totalThreads);
    long start = System.currentTimeMillis();
    for (int j = 0; j < numGroups; j++) {
        for (int i = 0; i < groupSize; i++) {
            Runnable thread = () -> {
                for (int k = 0; k < 1000; k++) {
                    clientGet.run();
                    clientPost.run();
                }
                completed2.countDown();
                //System.out.println("Thread completed");
            };
            new Thread(thread).start();
            //System.out.println("Thread started");
        }
        if (j != numGroups - 1) {
            Thread.sleep(delay * 1000);
        }
    }
}
```

NonRunnable Method

- Simple Get and Post Function
- ThreadPoolExecutor
- Scheduler

```
private static void sendGetRequest(int i) throws IOException { 2 usages

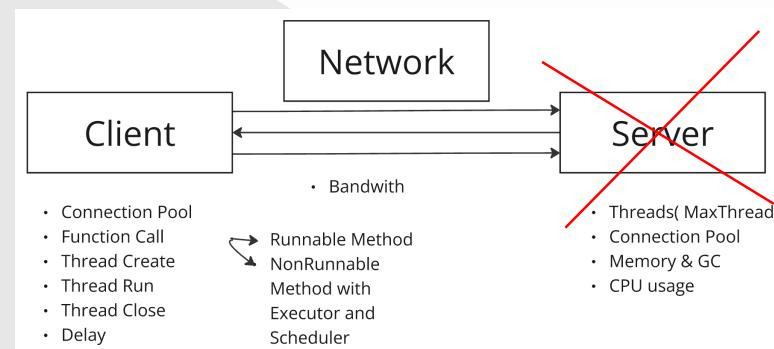
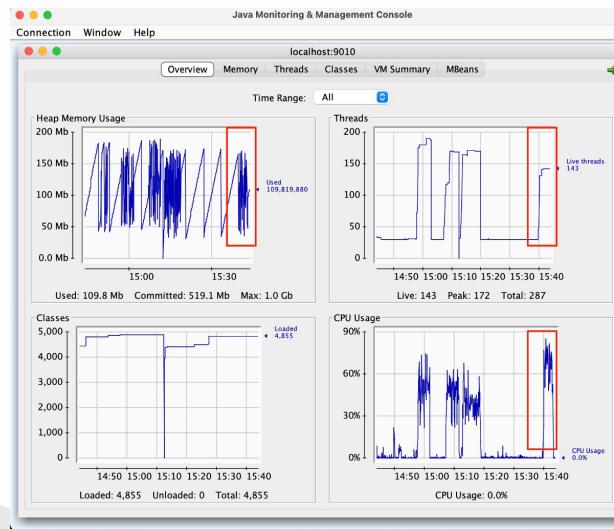
    long start = System.currentTimeMillis();
    HttpGet getMethod = new HttpGet(uri: IPAddrTomcat + "/" + i);
    try (CloseableHttpResponse response = client.execute(getMethod)) {
        EntityUtils.consume(response.getEntity());
    }
    long end = System.currentTimeMillis();
}
CountDownLatch latch2 = new CountDownLatch(threadGroupSize * numThreadGroups);
long start = System.currentTimeMillis();
for (int i = 0; i < numThreadGroups; i++) {
    int groupIndex = i;
    scheduler.schedule(() -> {
        for (int j = 0; j < threadGroupSize; j++) {
            int finalJ = j;
            executorService.submit(() -> {
                for (int k = 0; k < 1000; k++) {
                    try {
                        sendPostRequest();
                        sendGetRequest((groupIndex * threadGroupSize + finalJ));
                    } catch (IOException e) {
                        throw new RuntimeException(e);
                    }
                }
                latch2.countDown();
            });
        }
    }, delay: groupIndex * delay, TimeUnit.MILLISECONDS);
}
```

Experiments & Results

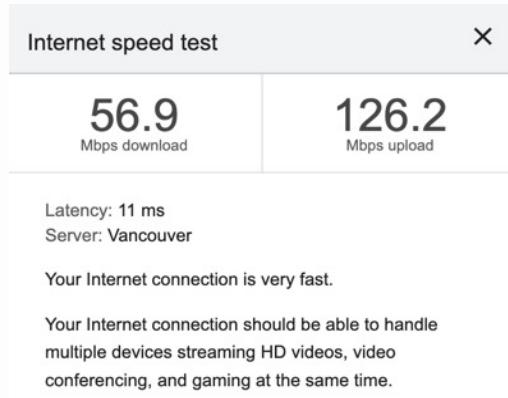
Biggest Problem – X Server

Time: 2025-02-12 09:44:12
Used: 136,516 kbytes
Committed: 506,944 kbytes
Max: 1,013,632 kbytes
GC time: 2.766 seconds on Copy (450 collections)
0.081 seconds on MarkSweepCompact (1 collections)

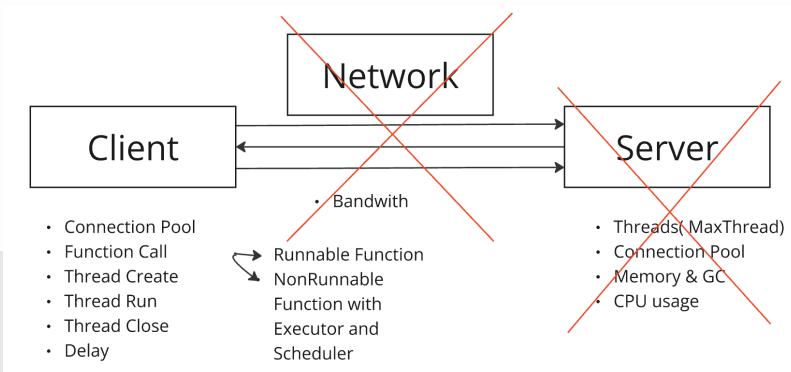
```
cp6      0  207 172.31.28.40:8080  208.98.212.98:23432 ESTABLISHED
cp6      0  0 172.31.28.40:8080  208.98.212.98:11907 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:51382 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:29608 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:5834 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:27109 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:48708 ESTABLISHED
cp6      0  0 172.31.28.40:8080  208.98.212.98:49893 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:27654 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:52310 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:16592 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:52143 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:28164 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:15015 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:29522 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:53808 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:21799 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:48719 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:49503 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:48478 ESTABLISHED
cp6      0  0 172.31.28.40:8080  208.98.212.98:18808 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:50996 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:30638 ESTABLISHED
cp6      0  207 172.31.28.40:8080  208.98.212.98:5479 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:58704 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:27059 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:52680 ESTABLISHED
cp6      0  0 172.31.28.40:8080  208.98.212.98:13012 ESTABLISHED
cp6      0  201 172.31.28.40:8080  208.98.212.98:48540 ESTABLISHED
```



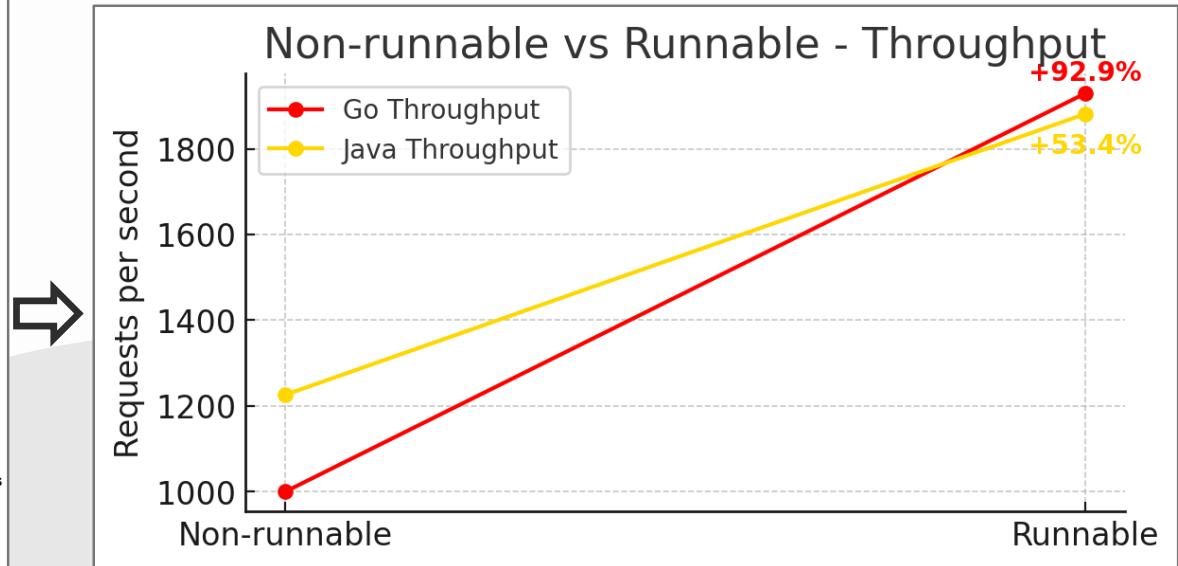
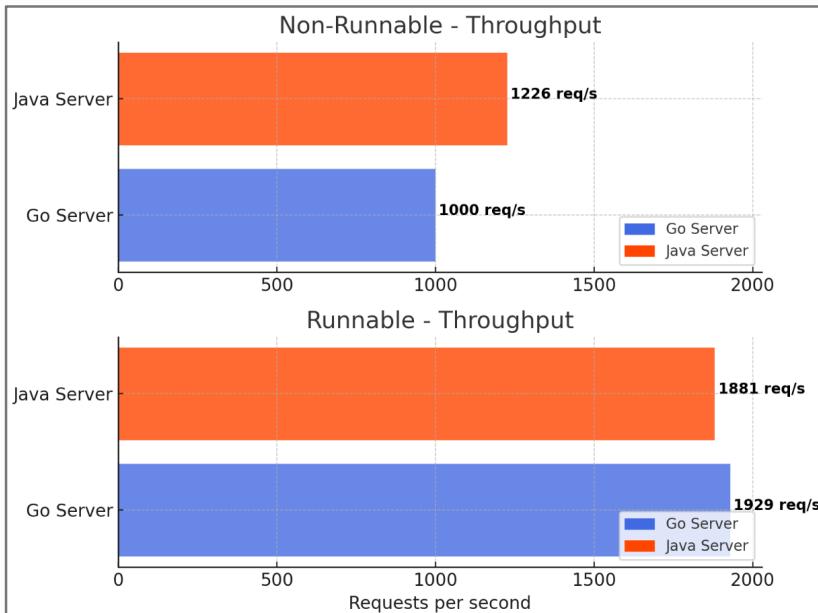
Biggest Problem – X Network



```
nuanxin@kindans-MacBook-Air ~ % traceroute 52.34.94.103
traceroute to 52.34.94.103 (52.34.94.103), 64 hops max, 40 byte packets
 1  10.247.192.1 (10.247.192.1)  5.013 ms  6.340 ms  7.387 ms
 2  * * *
 3  rc1st-be31-1.vc.shawcable.net (66.163.69.45)  4.847 ms  6.181 ms  6.082 ms
 4  rc2wh-tge0-13-0-2.vc.shawcable.net (66.163.69.97)  5.450 ms  5.753 ms *
 5  rd3st-tge0-11-0-0.vc.shawcable.net (66.163.69.90)  18.665 ms  7.298 ms  5.9
 5 ms
 6  * * *
 7  * 108.166.228.44 (108.166.228.44)  17.002 ms *
 8  * 108.166.228.44 (108.166.228.44)  19.157 ms  16.291 ms
 9  * * *
10  * * 108.166.228.63 (108.166.228.63)  21.684 ms
11  * * 108.166.228.83 (108.166.228.83)  18.014 ms
12  * * *
```

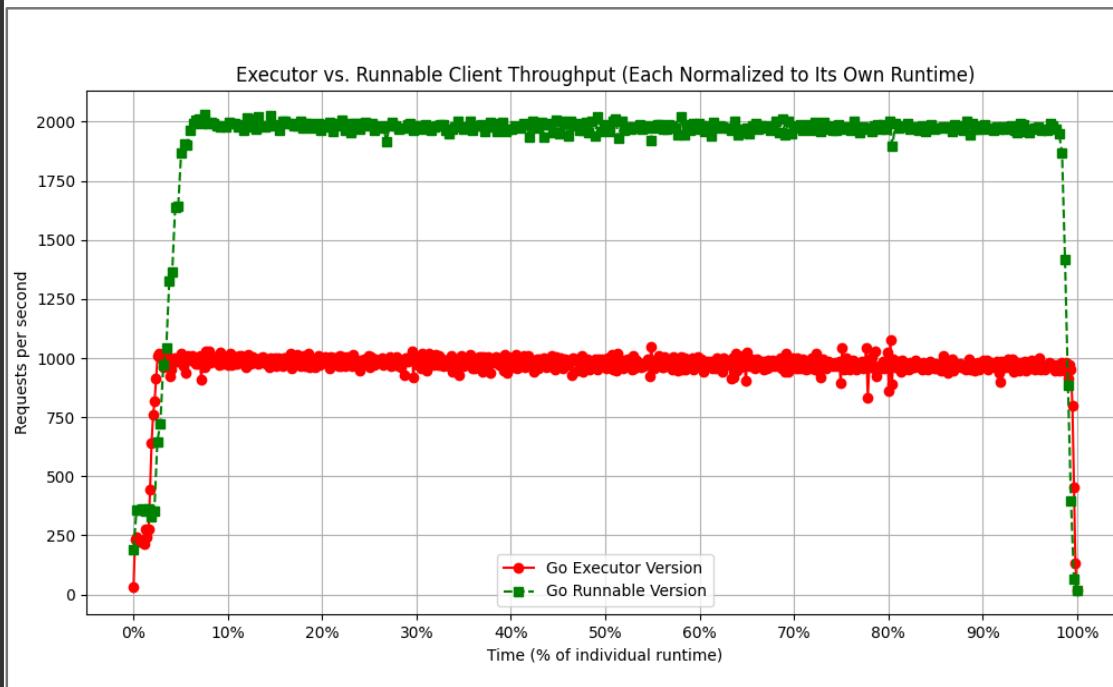


Runnable Increase Throughputs

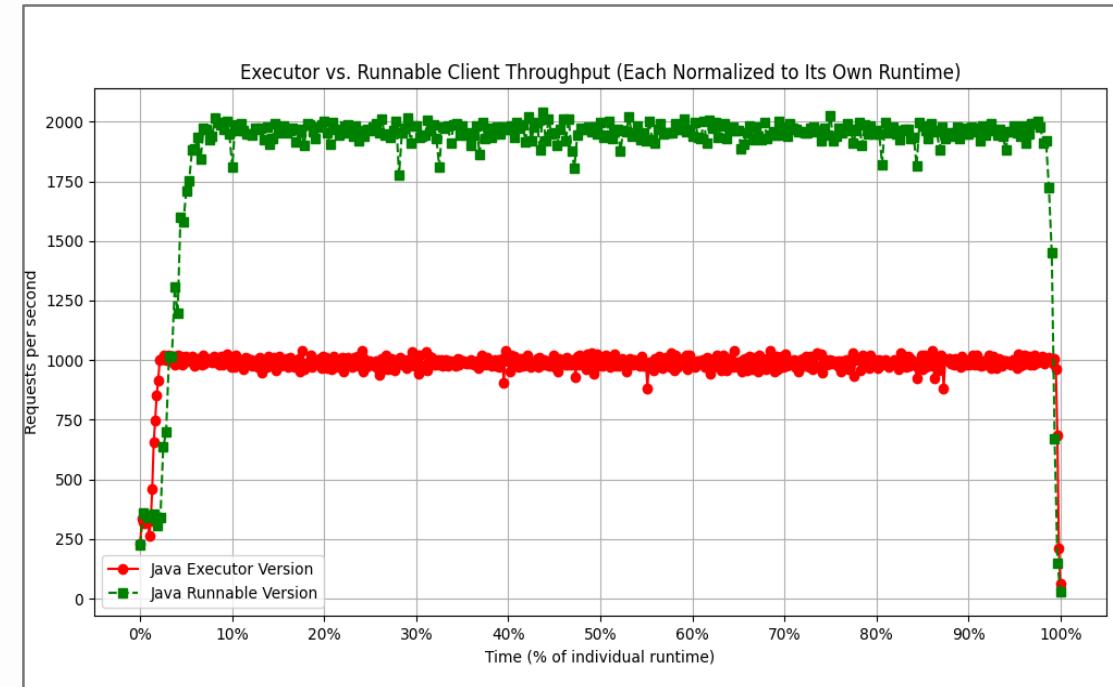


Throughput ↑ → Increased request handling capacity.

Runnable Increase Throughputs



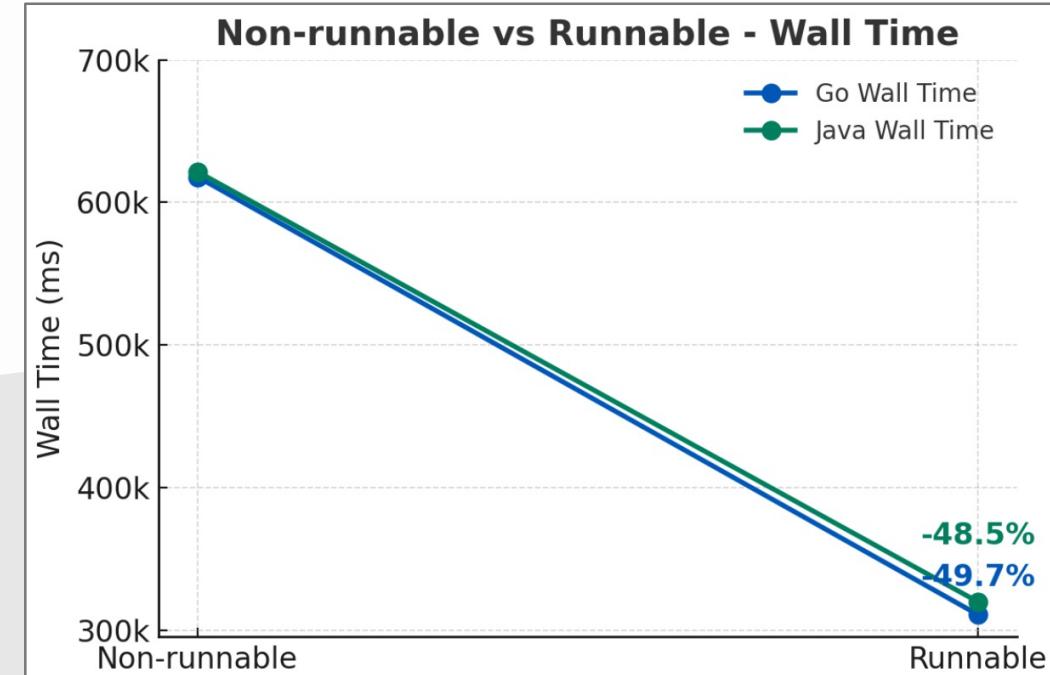
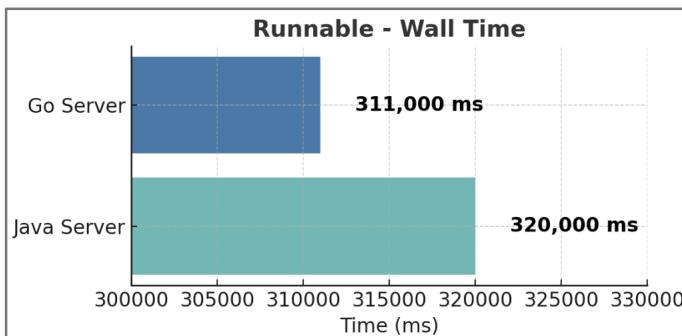
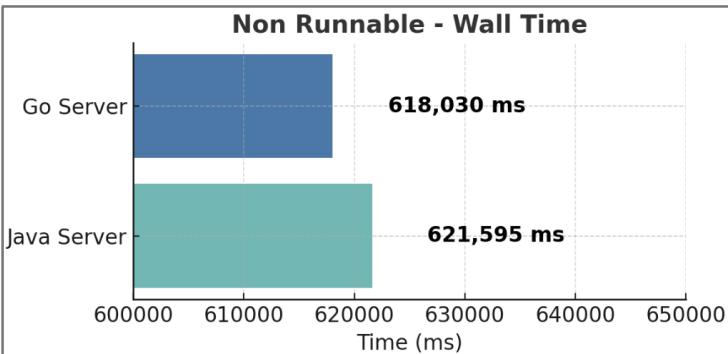
Go Server



Java Server

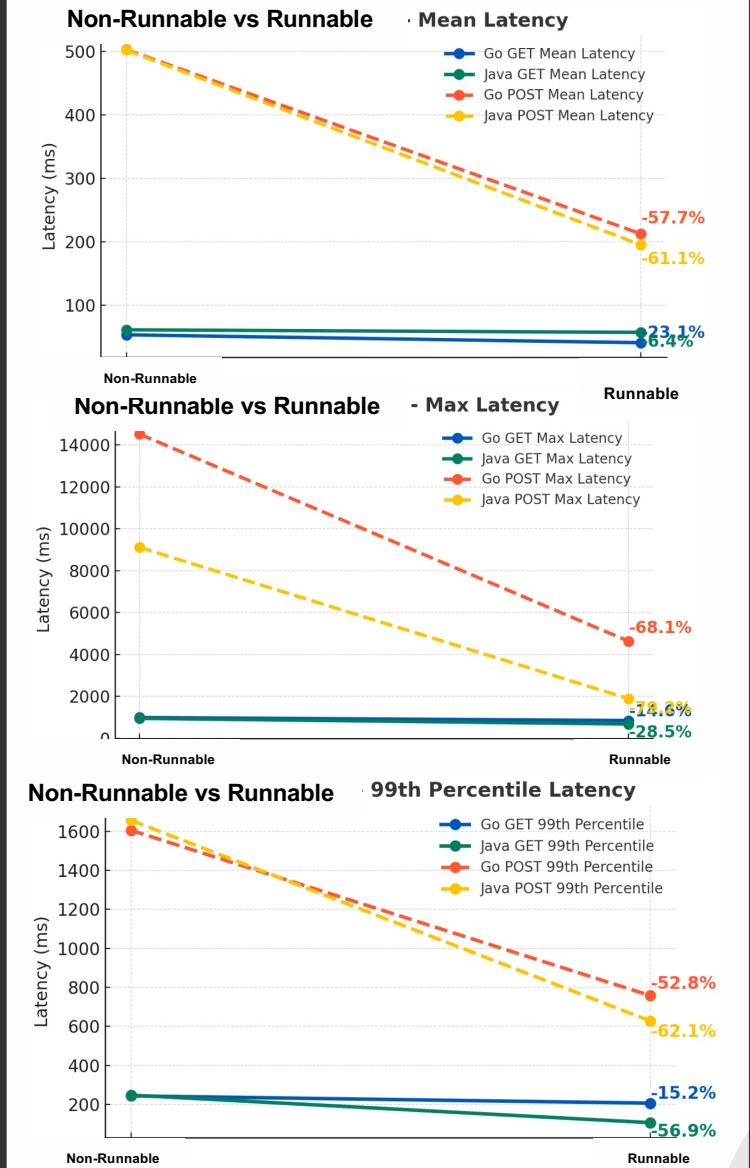
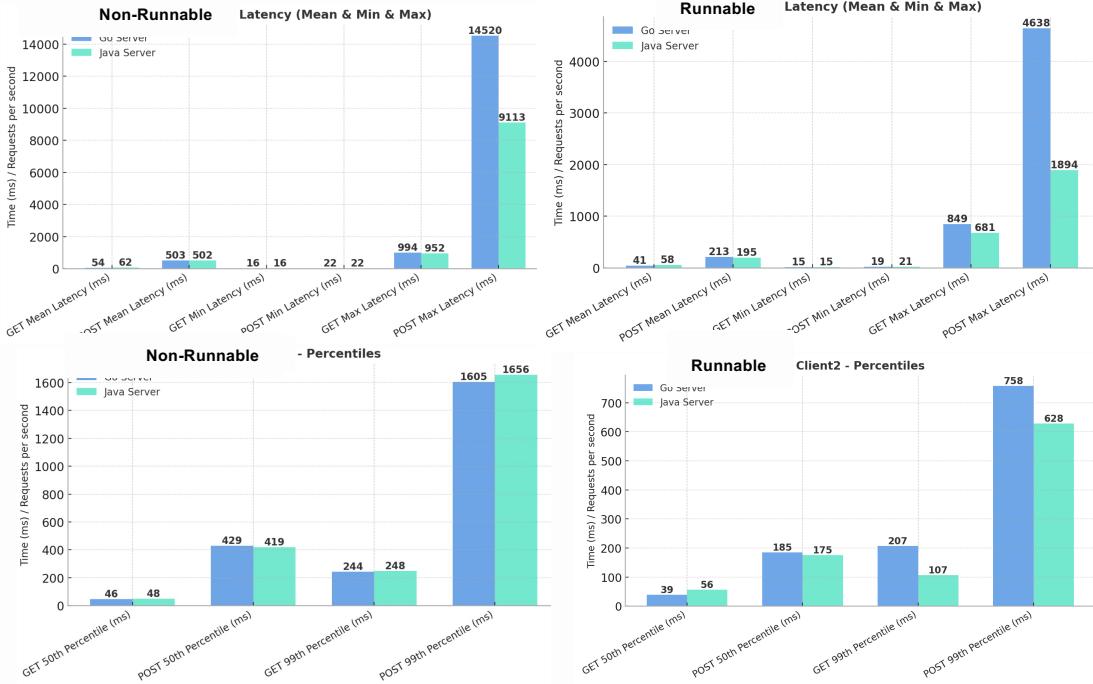
Runnable consistently achieves higher throughput

Runnable Reduce Wall Times



Wall Time ↓ → Faster execution and response times.

Runnable Reduce Latency



- **Mean Latency ↓:** Up to **61.1%**, with **POST requests improving the most.**
- **Max Latency ↓:** **68.1% (Go)** and **79.2% (Java)** reduction, minimizing long-tail delays.
- **99th Percentile ↓:** Improved stability, with up to **62.1%** fewer extreme slow requests.

Why Runnable

Client

- Connection Pool
- Function Call
- Thread Create
- Thread Run
- Thread Close
- Delay

Wall Time: 163560ms
Throughput: 1226 requests/second

Wall Time: 168463ms
Throughput: 1190 requests/second

getActiveCount()
getQueue().size()
getCompletedTaskCount()

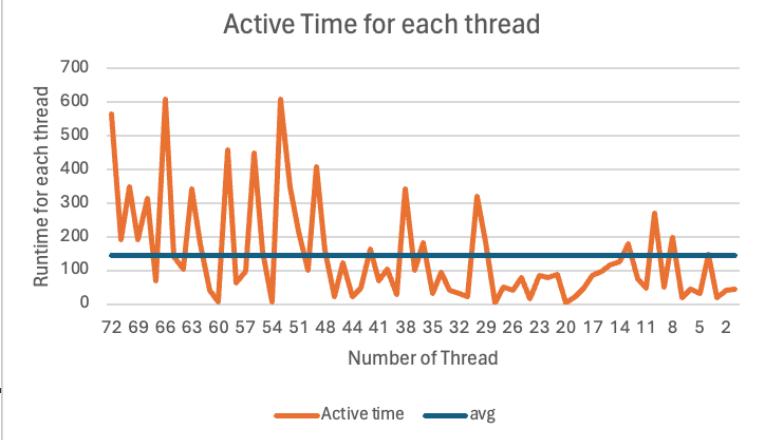
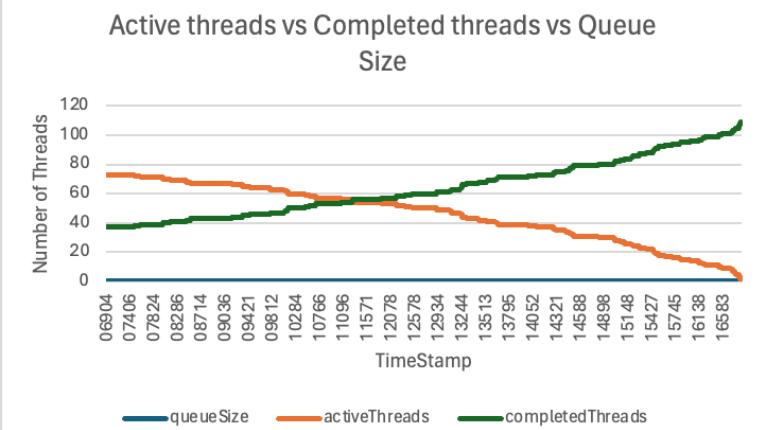
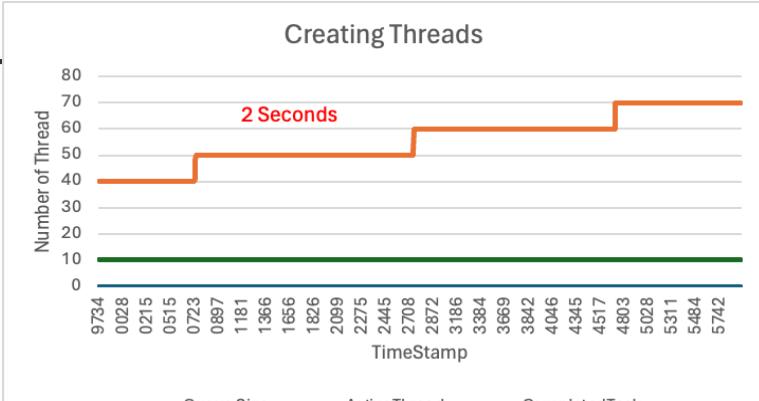
tcp6	0	201	172.31.28.40:8080	208.98.212.98:14638	ESTABLISHED
tcp6	0	207	172.31.28.40:8080	208.98.212.98:28850	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:19719	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:36708	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:35073	ESTABLISHED
tcp6	0	0	172.31.28.40:8080	208.98.212.98:45511	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:41689	ESTABLISHED
tcp6	0	207	172.31.28.40:8080	208.98.212.98:34710	ESTABLISHED
tcp6	0	207	172.31.28.40:8080	208.98.212.98:43824	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:26099	ESTABLISHED
tcp6	0	0	172.31.28.40:8080	208.98.212.98:17760	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:49679	ESTABLISHED
tcp6	0	201	172.31.28.40:8080	208.98.212.98:8072	ESTABLISHED

Runnable
Method

- NO Pool
- Runnable threads
- Thread Create instantly
- Thread Run instantly
- Thread Close automatically

NonRunnable
Method

- Fixed Pool > requirement Or CachedPool
- Simple Function
- Thread Create Managed by executor + Scheduler
- Thread Run Managed by executor + Scheduler
- Thread Close Manually

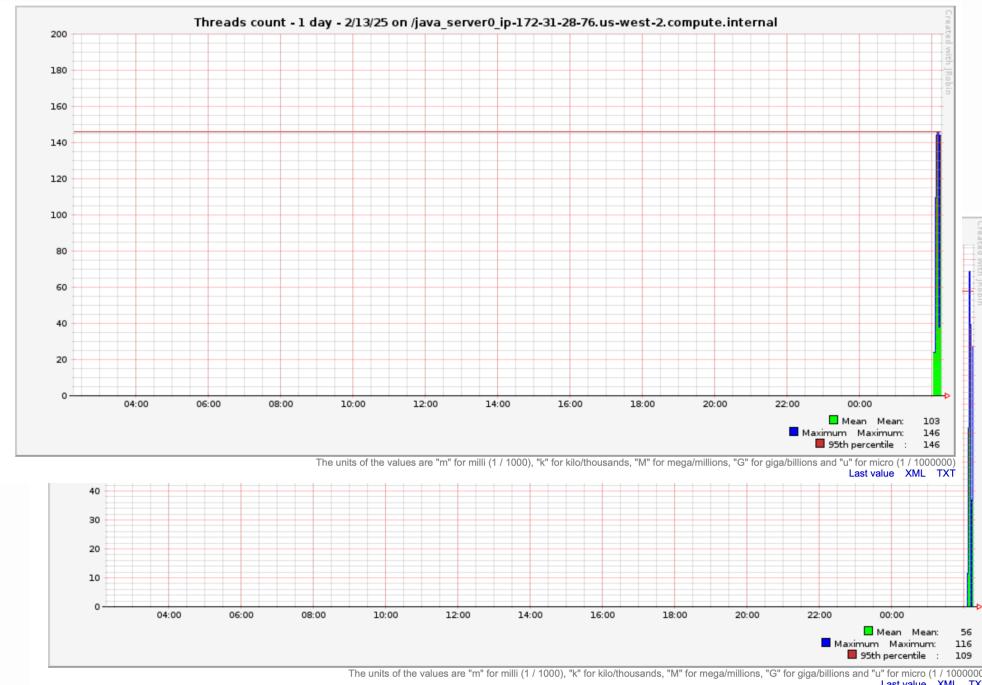


Chasing on Server

Scaling up server - How about launching another EC2?



Client: 300 threads



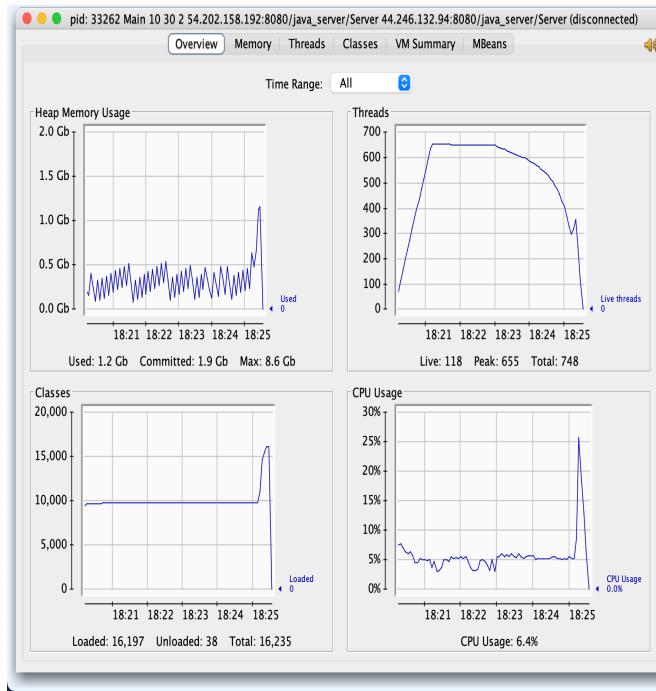
100+ threads / EC2

Time taken: 311s
Throughput: 1929 requests/s
SLF4J: Failed to load class "org.slf4j.
SLF4J: Defaulting to no-operation MDCAd:
SLF4J: See <http://www.slf4j.org/codes.html>
GET Mean Latency: 39.72816611295681
POST Mean Latency: 213.74947508305647
GET Min Latency: 15
POST Min Latency: 21
GET Max Latency: 508
POST Max Latency: 3011
GET 50th Percentile: 37.0
POST 50th Percentile: 177.0
GET 99th Percentile: 201.0
POST 99th Percentile: 804.0

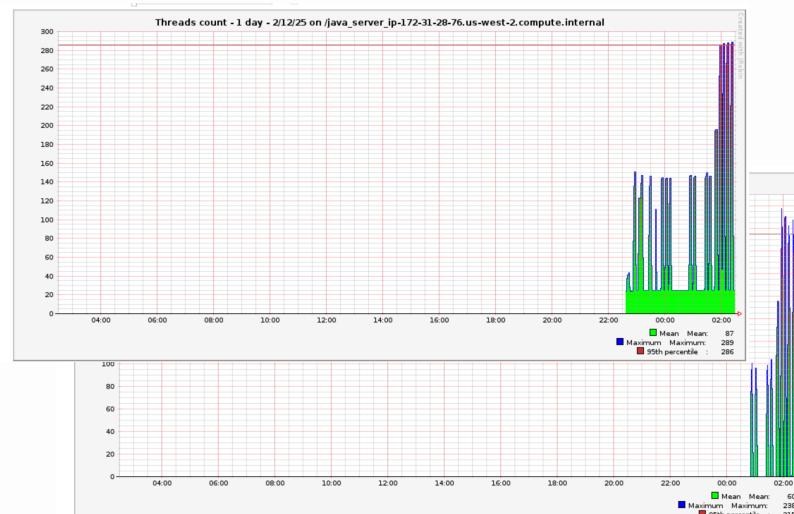
Throughput ~5%↑

Chasing on Server

More Threads from the Client!



Client: 600 threads



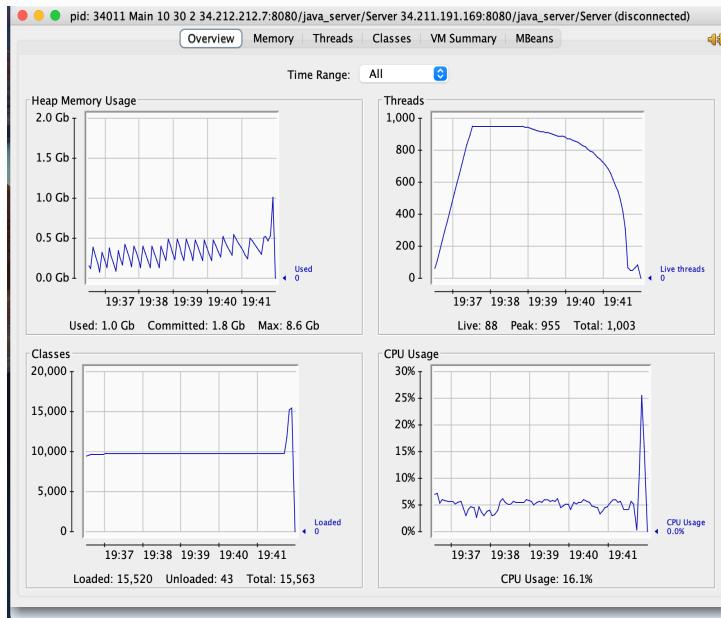
250+ Threads / EC2

Time taken: 295s
Throughput: 2033 requests/s
SLF4J: Failed to load class "org.slf4j.."
SLF4J: Defaulting to no-operation MDCAd
SLF4J: See <http://www.slf4j.org/codes.html#log4j2>
GET Mean Latency: 45.847362433754306
POST Mean Latency: 463.3617679738785
GET Min Latency: 15
POST Min Latency: 21
GET Max Latency: 1143
POST Max Latency: 15355
GET 50th Percentile: 38.0
POST 50th Percentile: 370.0
GET 99th Percentile: 214.0
POST 99th Percentile: 1663.0

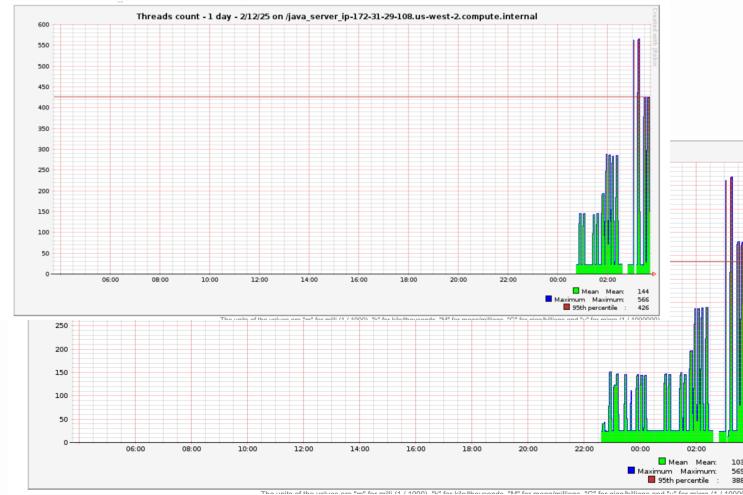
Throughput ~5% ↑

Chasing on Server

- Even More Threads!
- Throughput not improved
- POST Mean Latency Increased



Client 900 Threads



~450 Threads/ EC2

```
Time taken: 308s
Throughput: 1946 requests/s
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder"
SLF4J: Defaulting to no-operation MDCAdapter for category [org.slf4j]
SLF4J: See http://www.slf4j.org/codes.html#logback
GET Mean Latency: 46.7095942800133
POST Mean Latency: 698.6709078816095
GET Min Latency: 15
POST Min Latency: 21
GET Max Latency: 1651
POST Max Latency: 43764
GET 50th Percentile: 39.0
POST 50th Percentile: 551.0
GET 99th Percentile: 223.0
POST 99th Percentile: 2543.0
```

Throughput

Chasing on Network

This is an example image

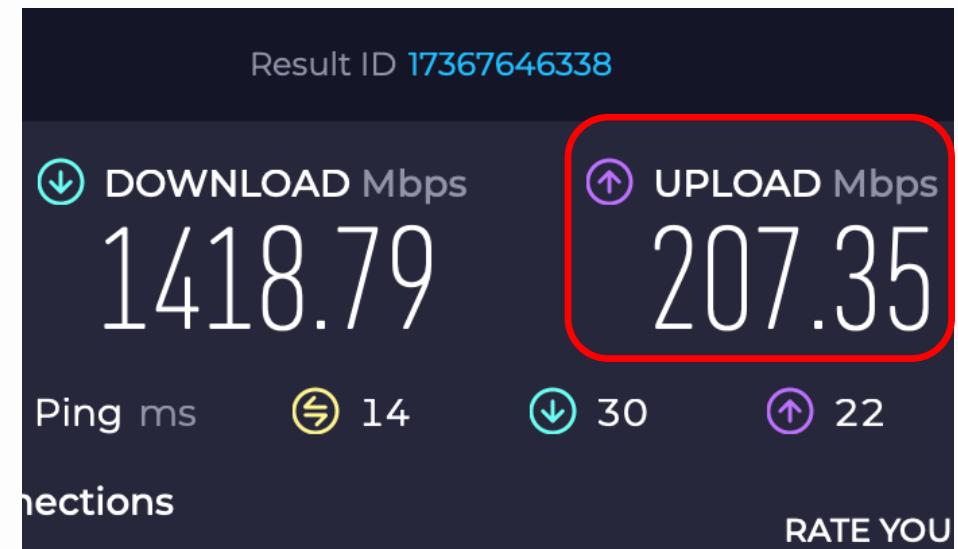


To use your own image, click
"Upload file" in the sidebar.

Images uploaded to Wikipedia must
be published under a free license.

25KB Each

$$25\text{KB} \times 1,000/\text{s} = 25\text{MB/s}$$

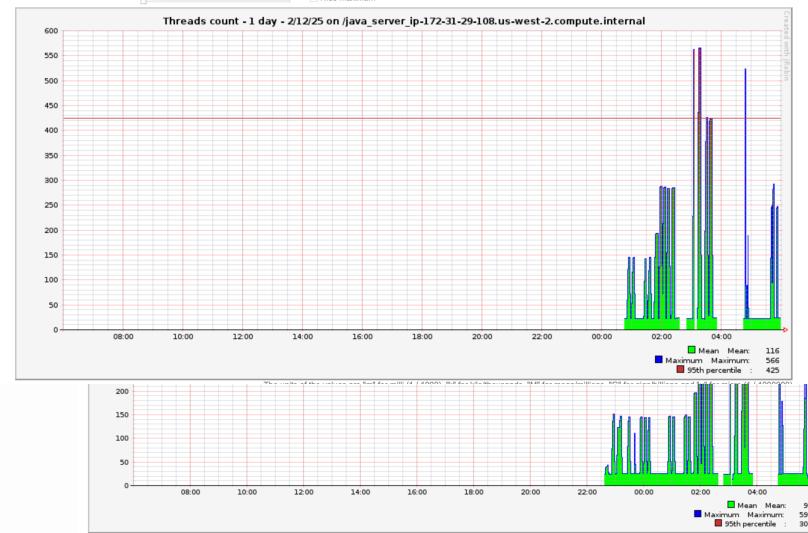


Chasing on Network

- Smaller Pic (2KB)! 
- Throughput is more than doubled!



Client 600 Threads



EC2 ~250 Threads each

Time taken: 129s
Throughput: 4651 requests/s
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder"
SLF4J: Defaulting to no-operation MDCAdap...
SLF4J: See <http://www.slf4j.org/codes.html>
GET Mean Latency: 43.72685382059801
POST Mean Latency: 90.64286046511629
GET Min Latency: 16
POST Min Latency: 20
GET Max Latency: 432
POST Max Latency: 859
GET 50th Percentile: 37.0
POST 50th Percentile: 75.0
GET 99th Percentile: 112.0
POST 99th Percentile: 316.0

Throughput Doubled!

Chasing the Bottleneck

- Next Bottleneck: File Descriptors! Too many Files Open!

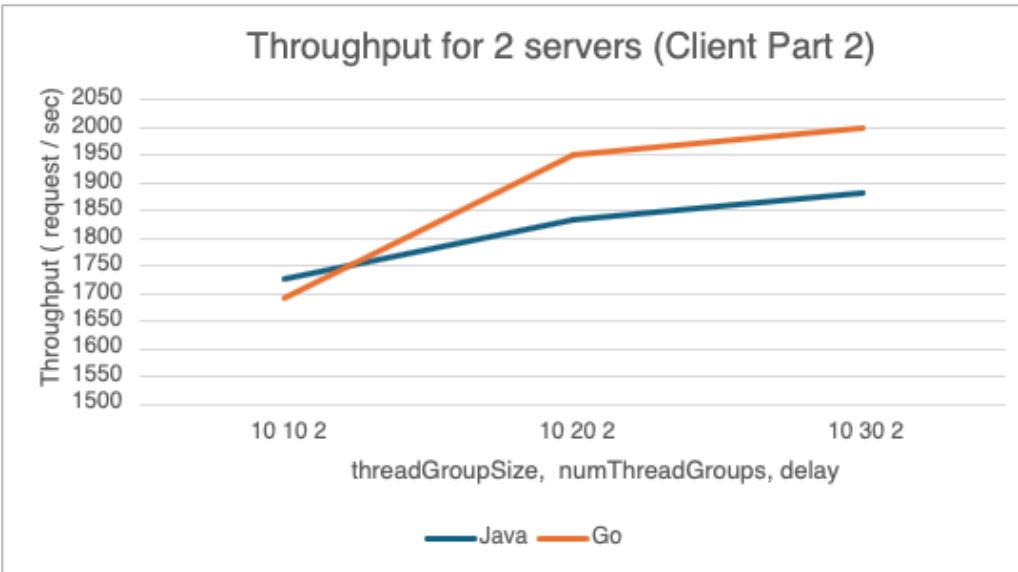
```
2/12/25, 4:53:56 AM | SEVERE: Servlet.service() for servlet [Server] in context with path [/java_server] threw exception
2/12/25, 4:53:56 AM | SEVERE: Servlet.service() for servlet [Server] in context with path [/java_server] threw exception
2/12/25, 4:53:56 AM | SEV org.apache.tomcat.util.http.fileupload.impl.IOFileUploadException: Processing of multipart/form-data request failed. /usr/share/apache-tomcat-9.0.80/work/Catalina/localhost/java_server/upload_76d43006_6430_4ee6_80f2_e4c3d37d2750_00872851.tmp (Too many open files)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.http.fileupload.FileUploadBase.parseRequest(FileUploadBase.java:321)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.connector.Request.parseParts(Request.java:283)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.connector.Request.getParts(Request.java:2739)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.connector.RequestFacade.getParts(RequestFacade.java:796)
2/12/25, 4:53:56 AM | SEV at javax.servlet.http.HttpServletWrapper.getParts(HttpServletWrapper.java:326)
2/12/25, 4:53:56 AM | SEV at Server doPost(Server.java:57)
2/12/25, 4:53:56 AM | SEV at javax.servlet.http.HttpServlet.service(HttpServlet.java:555)
2/12/25, 4:53:56 AM | SEV at javax.servlet.http.HttpServlet.service(HttpServletRequest.java:623)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.internalDoFilter(ApplicationFilterChain.java:209)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.doFilter(ApplicationFilterChain.java:153)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.websocket.server.WsFilter.doFilter(WsFilter.java:51)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.internalDoFilter(ApplicationFilterChain.java:178)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.doFilter(ApplicationFilterChain.java:153)
2/12/25, 4:53:56 AM | SEV at net.bull.javamelody.MonitoringFilter.doFilter(MonitoringFilter.java:239)
2/12/25, 4:53:56 AM | SEV at net.bull.javamelody.MonitoringFilter.doFilter(MonitoringFilter.java:215)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.internalDoFilter(ApplicationFilterChain.java:178)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.ApplicationFilterChain.doFilter(ApplicationFilterChain.java:153)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.StandardWrapperValve.invoke(StandardWrapperValve.java:168)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.StandardContextValve.invoke(StandardContextValve.java:90)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.authenticator.AuthenticatorBase.invoke(AuthenticatorBase.java:481)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.StandardHostValve.invoke(StandardHostValve.java:130)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.valves.ErrorReportValve.invoke(ErrorReportValve.java:93)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.valves.AbstractAccessLogValve.invoke(AbstractAccessLogValve.java:670)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.core.StandardEngineValve.invoke(StandardEngineValve.java:74)
2/12/25, 4:53:56 AM | SEV at org.apache.catalina.connector.CoyoteAdapter.service(CoyoteAdapter.java:343)
2/12/25, 4:53:56 AM | SEV at org.apache.coyote.http11.Http11Processor.service(Http11Processor.java:390)
2/12/25, 4:53:56 AM | SEV at org.apache.coyote.AbstractProcessorLight.process(AbstractProcessorLight.java:63)
2/12/25, 4:53:56 AM | SEV at org.apache.coyote.AbstractProtocol$ConnectionHandler.process(AbstractProtocol.java:926)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.net.NioEndpoint$SocketProcessor.doRun(NioEndpoint.java:1790)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.net.SocketProcessorBase.run(SocketProcessorBase.java:52)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.threads.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1191)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.threads.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:659)
2/12/25, 4:53:56 AM | SEV at org.apache.tomcat.util.threads.TaskThread$WrappingRunnable.run(TaskThread.java:61)
2/12/25, 4:53:56 AM | SEV at java.base/java.lang.Thread.run(Thread.java:829)
Caused by: java.io.FileNotFoundException: /usr/share/apache-tomcat-9.0.80/work/Catalina/localhost/java_server/upload_76d43006_6430_4ee6_80f2_e4c3d37d2750_00872851.tmp (Too many open files)
at java.base/java.io.FileOutputStream.open0(Native Method)
at java.base/java.io.FileOutputStream.open(FileOutputStream.java:298)
at java.base/java.io.FileOutputStream.<init>(FileOutputStream.java:237)
at java.base/java.io.FileOutputStream.<init>(FileOutputStream.java:187)
at org.apache.tomcat.util.http.fileupload.DeferredFileOutputStream.thresholdReached(DeferredFileOutputStream.java:151)
at org.apache.tomcat.util.http.fileupload.ThresholdingOutputStream.checkThreshold(ThresholdingOutputStream.java:200)
at org.apache.tomcat.util.http.fileupload.ThresholdingOutputStream.write(ThresholdingOutputStream.java:126)
at org.apache.tomcat.util.http.fileupload.util.Streams.copy(Streams.java:104)
at org.apache.tomcat.util.http.fileupload.FileUploadBase.parseRequest(FileUploadBase.java:317)
... 33 more
2/12/25, 4:53:56 AM | SEVERE: Servlet.service() for servlet [Server] in context with path [/java_server] threw exception
```

Future Works

Future Experiments with Clients

- Executor introduces overhead?
- Eager Thread Creation vs. Lazy Thread Creation
- Older threads delaying newer threads?
-

Future Experiments with Server



Go Outperforms Java

Goroutine	Java thread
Managed by go runtime	Managed by operating system, kernel
Can be resized (like dynamic array in C!)	Fixed size
Initial stack size ~2KB	~1MB by default
Low overhead :)	High overhead :(

Future Experiments with Server - Scaling



Kubernetes

automates the deployment,
scaling, and management
of containerized
applications



Auto Scaling Group

automatically adjust the
number of instances
based on traffic demand

Thanks!